

Node.js Fundamentals

Dave van Soest

January 13, 2015

Technology

Asynchronous programming

Examples

API

Modules

Packages

Technologies

- V8 (Javascript engine created by Google)

Technologies

- V8 (Javascript engine created by Google)
- C++ (Libraries)

Asynchronous programming

Single-threaded process

The Node application is executed as a **single-threaded process**.

Asynchronous programming

Single-threaded process

The Node application is executed as a **single-threaded process**.

Consequences

Asynchronous programming

Single-threaded process

The Node application is executed as a **single-threaded process**.

Consequences

- The application can only be busy doing one thing at a time.

Asynchronous programming

Single-threaded process

The Node application is executed as a **single-threaded process**.

Consequences

- The application can only be busy doing one thing at a time.
- Long-running calculations will block other work.

Asynchronous programming

Single-threaded process

The Node application is executed as a **single-threaded process**.

Consequences

- The application can only be busy doing one thing at a time.
- Long-running calculations will block other work.
- Blocking system calls will pause the application.

Asynchronous programming

Non-blocking I/O

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.
- Non-blocking system calls return immediately, leaving the result of the operation to be fetched with a later system call.

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.
- Non-blocking system calls return immediately, leaving the result of the operation to be fetched with a later system call.

Simplified example

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.
- Non-blocking system calls return immediately, leaving the result of the operation to be fetched with a later system call.

Simplified example

- Start reading a block of data from a file descriptor (system data resource) using non-blocking I/O.

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.
- Non-blocking system calls return immediately, leaving the result of the operation to be fetched with a later system call.

Simplified example

- Start reading a block of data from a file descriptor (system data resource) using non-blocking I/O.
- Loop:

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.
- Non-blocking system calls return immediately, leaving the result of the operation to be fetched with a later system call.

Simplified example

- Start reading a block of data from a file descriptor (system data resource) using non-blocking I/O.
- Loop:
 - Check whether read operation is completed.

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.
- Non-blocking system calls return immediately, leaving the result of the operation to be fetched with a later system call.

Simplified example

- Start reading a block of data from a file descriptor (system data resource) using non-blocking I/O.
- Loop:
 - Check whether read operation is completed.
 - If read data available:

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.
- Non-blocking system calls return immediately, leaving the result of the operation to be fetched with a later system call.

Simplified example

- Start reading a block of data from a file descriptor (system data resource) using non-blocking I/O.
- Loop:
 - Check whether read operation is completed.
 - If read data available:
 - Do something with the data.

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.
- Non-blocking system calls return immediately, leaving the result of the operation to be fetched with a later system call.

Simplified example

- Start reading a block of data from a file descriptor (system data resource) using non-blocking I/O.
- Loop:
 - Check whether read operation is completed.
 - If read data available:
 - Do something with the data.
 - If read error:

Asynchronous programming

Non-blocking I/O

- Operations using peripheral hardware (disks, network interfaces, timers, ...) take a relatively long time to complete.
- Non-blocking system calls return immediately, leaving the result of the operation to be fetched with a later system call.

Simplified example

- Start reading a block of data from a file descriptor (system data resource) using non-blocking I/O.
- Loop:
 - Check whether read operation is completed.
 - If read data available:
 - Do something with the data.
 - If read error:
 - Display error message.

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.
 - If new incoming connection:

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.
 - If new incoming connection:
 - Start receiving data from incoming connection (non-blocking).

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.
 - If new incoming connection:
 - Start receiving data from incoming connection (non-blocking).
 - If data received from connection:

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.
 - If new incoming connection:
 - Start receiving data from incoming connection (non-blocking).
 - If data received from connection:
 - Process data and generate response.

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.
 - If new incoming connection:
 - Start receiving data from incoming connection (non-blocking).
 - If data received from connection:
 - Process data and generate response.
 - Start sending response data to connection (non-blocking).

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.
 - If new incoming connection:
 - Start receiving data from incoming connection (non-blocking).
 - If data received from connection:
 - Process data and generate response.
 - Start sending response data to connection (non-blocking).
 - If response data is sent:

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.
 - If new incoming connection:
 - Start receiving data from incoming connection (non-blocking).
 - If data received from connection:
 - Process data and generate response.
 - Start sending response data to connection (non-blocking).
 - If response data is sent:
 - Close the connection.

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.
 - If new incoming connection:
 - Start receiving data from incoming connection (non-blocking).
 - If data received from connection:
 - Process data and generate response.
 - Start sending response data to connection (non-blocking).
 - If response data is sent:
 - Close the connection.

Callback functions

Asynchronous programming

Non-blocking I/O: Multiple resources

Example: Multiple data resources

- Listen for incoming network connections (non-blocking).
- Loop:
 - Check for events.
 - If new incoming connection:
 - Start receiving data from incoming connection (non-blocking).
 - If data received from connection:
 - Process data and generate response.
 - Start sending response data to connection (non-blocking).
 - If response data is sent:
 - Close the connection.

Callback functions

For each combination of *[resource, event-type]* the application registers a **callback function** to be executed when the event occurs.

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

- Continuous checking (a.k.a. polling) by the application would use a lot of system resources.

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

- Continuous checking (a.k.a. polling) by the application would use a lot of system resources.
- Solution: The *select*-family of system calls (blocking).

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

- Continuous checking (a.k.a. polling) by the application would use a lot of system resources.
- Solution: The *select*-family of system calls (blocking).
 - Monitor multiple system resources for events at the same time.

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

- Continuous checking (a.k.a. polling) by the application would use a lot of system resources.
- Solution: The *select*-family of system calls (blocking).
 - Monitor multiple system resources for events at the same time.
 - Return when an event occurs, a timer expires or a signal is received.

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

- Continuous checking (a.k.a. polling) by the application would use a lot of system resources.
- Solution: The *select*-family of system calls (blocking).
 - Monitor multiple system resources for events at the same time.
 - Return when an event occurs, a timer expires or a signal is received.
 - Uses minimal amount of system resources. Application itself is idle.

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

- Continuous checking (a.k.a. polling) by the application would use a lot of system resources.
- Solution: The *select*-family of system calls (blocking).
 - Monitor multiple system resources for events at the same time.
 - Return when an event occurs, a timer expires or a signal is received.
 - Uses minimal amount of system resources. Application itself is idle.
- Leverages:

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

- Continuous checking (a.k.a. polling) by the application would use a lot of system resources.
- Solution: The *select*-family of system calls (blocking).
 - Monitor multiple system resources for events at the same time.
 - Return when an event occurs, a timer expires or a signal is received.
 - Uses minimal amount of system resources. Application itself is idle.
- Leverages:
 - Software interrupts

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

- Continuous checking (a.k.a. polling) by the application would use a lot of system resources.
- Solution: The *select*-family of system calls (blocking).
 - Monitor multiple system resources for events at the same time.
 - Return when an event occurs, a timer expires or a signal is received.
 - Uses minimal amount of system resources. Application itself is idle.
- Leverages:
 - Software interrupts
 - Hardware interrupts

Asynchronous programming

Non-blocking I/O: Optimizing performance

Optimizing performance

- Continuous checking (a.k.a. polling) by the application would use a lot of system resources.
- Solution: The *select*-family of system calls (blocking).
 - Monitor multiple system resources for events at the same time.
 - Return when an event occurs, a timer expires or a signal is received.
 - Uses minimal amount of system resources. Application itself is idle.
- Leverages:
 - Software interrupts
 - Hardware interrupts
 - Direct Memory Access (D.M.A.)

Asynchronous programming

Advantages

Advantages

Asynchronous programming

Advantages

Advantages

- Optimizing system resource usage:

Asynchronous programming

Advantages

Advantages

- Optimizing system resource usage:
 - Running operations on multiple data resources in parallel instead of sequential.

Asynchronous programming

Advantages

Advantages

- Optimizing system resource usage:
 - Running operations on multiple data resources in parallel instead of sequential.
 - Reduces relatively expensive process-switching and thread-switching.

Asynchronous programming

Advantages

Advantages

- Optimizing system resource usage:
 - Running operations on multiple data resources in parallel instead of sequential.
 - Reduces relatively expensive process-switching and thread-switching.
- No worries about thread-safety.

Asynchronous programming

Javascript

Javascript is perfectly suited for asynchronous programming.

Asynchronous programming

Javascript

Javascript is perfectly suited for asynchronous programming.

- Javascript engines have a built-in event-loop.

Asynchronous programming

Javascript

Javascript is perfectly suited for asynchronous programming.

- Javascript engines have a built-in event-loop.
- First-class functions are used as callbacks.

Asynchronous programming

Javascript

Javascript is perfectly suited for asynchronous programming.

- Javascript engines have a built-in event-loop.
- First-class functions are used as callbacks.
- Anonymous functions allow for in-line callbacks.

Asynchronous programming

Javascript

Javascript is perfectly suited for asynchronous programming.

- Javascript engines have a built-in event-loop.
- First-class functions are used as callbacks.
- Anonymous functions allow for in-line callbacks.
- Automatic binding of a function to its outside context (closure) allows for easy access of program state inside a callback.

Simple example

HTTP-server example from Node.js homepage

```
1  var http = require("http");
2
3  var httpServer = http.createServer(
4    function (req, res) {
5      res.writeHead(200, {"Content-Type": "text/plain"});
6      res.end("Hello␣World␣\n");
7    }
8  );
9  httpServer.listen(1337, "127.0.0.1");
10
11 console.log("Server␣running␣at␣http://127.0.0.1:1337/");
```

Extended example

Simple file server

```
1  var http = require("http");
2  var fs = require("fs");
3
4  var httpServer = http.createServer(
5    function (req, res) {
6      fs.readFile(req.url, function(err, data) {
7        res.writeHead(err ? 404 : 200, {"Content-Type": "text/plain"});
8        res.end(err ? "Not found.\n" : data);
9      });
10   }
11 );
12 httpServer.listen(1337, "127.0.0.1");
```

Example using streams and event emitters

Streaming file server

```
1 var http = require("http");
2 var fs = require("fs");
3
4 var httpServer = http.createServer(
5   function (req, res) {
6     var stream = fs.createReadStream(req.url);
7
8     stream.on("open", function() {
9       res.writeHead(200, {"Content-Type": "text/plain"});
10      stream.pipe(res);
11    });
12
13    stream.on("error", function(err) {
14      res.writeHead(400, {"Content-Type": "text/plain"});
15      res.end(err);
16    });
17  }
18 );
19 httpServer.listen(1337, "127.0.0.1");
```


API

Selection from the Node.js built-in API:

API

Selection from the Node.js built-in API:

- **http**: HTTP-server and client.

API

Selection from the Node.js built-in API:

- http: HTTP-server and client.
- **net**: TCP server and client.

API

Selection from the Node.js built-in API:

- http: HTTP-server and client.
- net: TCP server and client.
- fs: Filesystem operations (reading, writing, listing, ...).

API

Selection from the Node.js built-in API:

- http: HTTP-server and client.
- net: TCP server and client.
- fs: Filesystem operations (reading, writing, listing, ...).
 - Both non-blocking and blocking operations!

API

Selection from the Node.js built-in API:

- http: HTTP-server and client.
- net: TCP server and client.
- fs: Filesystem operations (reading, writing, listing, ...).
 - Both non-blocking and blocking operations!
- path: File path handling.

API

Selection from the Node.js built-in API:

- http: HTTP-server and client.
- net: TCP server and client.
- fs: Filesystem operations (reading, writing, listing, ...).
 - Both non-blocking and blocking operations!
- path: File path handling.
- **readline**: Create a text-based user interface.

API

Selection from the Node.js built-in API:

- http: HTTP-server and client.
- net: TCP server and client.
- fs: Filesystem operations (reading, writing, listing, ...).
 - Both non-blocking and blocking operations!
- path: File path handling.
- readline: Create a text-based user interface.
- url, querystring: URL parsing and generating.

API

Selection from the Node.js built-in API:

- http: HTTP-server and client.
- net: TCP server and client.
- fs: Filesystem operations (reading, writing, listing, ...).
 - Both non-blocking and blocking operations!
- path: File path handling.
- readline: Create a text-based user interface.
- url, querystring: URL parsing and generating.
- **globals**: global, process, console, module, ...

API

Selection from the Node.js built-in API:

- http: HTTP-server and client.
- net: TCP server and client.
- fs: Filesystem operations (reading, writing, listing, ...).
 - Both non-blocking and blocking operations!
- path: File path handling.
- readline: Create a text-based user interface.
- url, querystring: URL parsing and generating.
- globals: global, process, console, module, ...
- ...

Modules

Using a module

Using a module

Modules

Using a module

Using a module

- Importing a module:

```
var xyz = require("xyz");  
(short-hand for module.require)
```

Modules

Using a module

Using a module

- Importing a module:

```
var xyz = require("xyz");  
(short-hand for module.require)
```

- Module is searched for in local `node_modules` directory and in the system's modules directory.

Modules

Using a module

Using a module

- Importing a module:

```
var xyz = require("xyz");
```

(short-hand for `module.require`)

- Module is searched for in local `node_modules` directory and in the system's modules directory.
- Importing a local module:

```
var abc = require("./abc");
```

Modules

Using a module

Using a module

- Importing a module:

```
var xyz = require("xyz");
```

(short-hand for `module.require`)

- Module is searched for in local `node_modules` directory and in the system's modules directory.
- Importing a local module:

```
var abc = require("./abc");
```

 - Module is searched for in local directory.

Modules

Module initialization

Example module

```
1 // Initialization code:
2 console.log("Hello␣ABC.");
3
4 // Exporting functionality:
5 exports.doAbc = function() {
6   console.log("Performing␣ABC.");
7 };
```


Modules

Module initialization

Example module

```
1 // Initialization code:
2 console.log("Hello␣ABC.");
3
4 // Exporting functionality:
5 exports.doAbc = function() {
6   console.log("Performing␣ABC.");
7 };
```

Module initialization

Modules

Module initialization

Example module

```
1 // Initialization code:
2 console.log("Hello␣ABC.");
3
4 // Exporting functionality:
5 exports.doAbc = function() {
6   console.log("Performing␣ABC.");
7 };
```

Module initialization

- A module is initialized upon the first import.

Modules

Module initialization

Example module

```
1 // Initialization code:
2 console.log("Hello␣ABC.");
3
4 // Exporting functionality:
5 exports.doAbc = function() {
6   console.log("Performing␣ABC.");
7 };
```

Module initialization

- A module is initialized upon the first import.
- Each module is only **initialized once** within the program.

Modules

Creating a module

Creating a module

Modules

Creating a module

Creating a module

- A local module can be ...

Modules

Creating a module

Creating a module

- A local module can be ...
 - a single file `abc.js` or a directory

Modules

Creating a module

Creating a module

- A local module can be ...
 - a single file `abc.js` or a directory
 - a directory `abc` with a file called `index.js` as the module's main entry point

Modules

Creating a module

Creating a module

- A local module can be ...
 - a single file `abc.js` or a directory
 - a directory `abc` with a file called `index.js` as the module's main entry point
 - a directory `abc` with a package definition file `package.json`

Modules

Creating a module

Creating a module

- A local module can be ...
 - a single file `abc.js` or a directory
 - a directory `abc` with a file called `index.js` as the module's main entry point
 - a directory `abc` with a package definition file `package.json`
 - This type of module is called a **package**.

Modules

Creating a module

Creating a module

- A local module can be ...
 - a single file `abc.js` or a directory
 - a directory `abc` with a file called `index.js` as the module's main entry point
 - a directory `abc` with a package definition file `package.json`
 - This type of module is called a package.
- A module can import other modules.

Modules

Creating a module

Creating a module

- A local module can be ...
 - a single file `abc.js` or a directory
 - a directory `abc` with a file called `index.js` as the module's main entry point
 - a directory `abc` with a package definition file `package.json`
 - This type of module is called a package.
- A module can import other modules.
- A module exports data and functionality through the `exports` object (short-hand for `module.exports`)

Modules

Local module usage example

./abc.js

```
1 // Initialization code:
2 console.log("Hello␣ABC.");
3
4 // Exporting functionality:
5 exports.doAbc = function() {
6     console.log("Performing␣ABC.");
7 };
```

Modules

Local module usage example

./abc.js

```
1 // Initialization code:
2 console.log("Hello␣ABC.");
3
4 // Exporting functionality:
5 exports.doAbc = function() {
6     console.log("Performing␣ABC.");
7 };
```

./app.js

```
1 var abc = require("./abc");
2
3 abc.doAbc();
```

Packages

Package definition

File: package.json in module directory

```
1 {  
2   "name": "abc",  
3   "version": "1.0.0",  
4   "description": "A, B and C",  
5   "main": "index.js",  
6   "dependencies": {  
7     "underscore": "~1.6.0"  
8     "xyz": "1.2.3"  
9   }  
10   ...  
11 }
```

Packages

Node Package Manager (NPM)

Packages

Node Package Manager (NPM)

- Installing packages locally (in `node_modules` directory) or globally.

Packages

Node Package Manager (NPM)

- Installing packages locally (in `node_modules` directory) or globally.
- Updating packages.

Packages

Node Package Manager (NPM)

- Installing packages locally (in `node_modules` directory) or globally.
- Updating packages.
- Package version management.

Packages

Node Package Manager (NPM)

- Installing packages locally (in `node_modules` directory) or globally.
- Updating packages.
- Package version management.
- Publishing packages.